

## 1. Problem Definition (6 points)

**Hypothetical AI Problem:** *Detecting biased language in online job postings.*

### Objectives:

1. Identify and flag gender-coded or discriminatory phrases.
2. Suggest neutral language alternatives in real time.
3. Improve diversity in applicant pools over time.

### Stakeholders:

- Recruitment teams
- Job seekers from underrepresented groups

**KPI:** *Reduction in flagged biased phrases per 100 job postings over time.*

## 2. Data Collection & Preprocessing (8 points)

### Data Sources:

1. Public job listing datasets (e.g., Kaggle’s job descriptions dataset)
2. Web-scraped job ads from company websites or job boards (e.g., Indeed, LinkedIn)

**Potential Bias:** Training data may already reflect historical biases (e.g., overuse of “aggressive,” “dominant,” “rockstar”), which could reinforce the status quo if not corrected.

### Preprocessing Steps:

- Text cleaning (punctuation removal, lowercasing)
- Tokenization and lemmatization
- Handling class imbalance (e.g., SMOTE or stratified sampling)

## 3. Model Development (8 points)

**Chosen Model:** *Fine-tuned BERT (Bidirectional Encoder Representations from Transformers)*

**Justification:** BERT is highly effective for NLP tasks and captures contextual meaning in language—essential for identifying subtle bias.

### Data Split Strategy:

- 70% training
- 15% validation
- 15% testing Stratified sampling ensures representation of different bias types.

### Hyperparameters to Tune:

1. Learning rate – to optimize training stability and convergence.
2. Maximum sequence length – to control context depth for longer job descriptions.

#### 4. Evaluation & Deployment (8 points)

##### Evaluation Metrics:

- *F1-Score*: Balances precision and recall—critical in detecting rare biased phrases.
- *False Positive Rate*: Ensures model doesn't over-flag neutral language.

**Concept Drift:** Occurs when the statistical properties of the data change over time (e.g., new slang, shifting definitions of inclusivity). **Monitoring Solution:** Use active learning or continuous evaluation pipelines with periodic re-training using recent data.

**Technical Deployment Challenge:** Scalability – processing large volumes of job ads in real time may require distributed computing or efficient model serving via APIs like FastAPI + TorchServe.